

Tutorials

Build AI Programs with DSPy

Managing Conversation History

Building AI Agents with DSPy

Building AI Applications by Customizing DSPy Modules

Retrieval-Augmented Generation (RAG)

Building RAG as Agent

Entity Extraction

Classification

Multi-Hop RAG

Privacy-Conscious Delegation

Program Of Thought

Image Generation Prompt Iteration

Audio

Optimize AI Programs with DSPy

Math Reasoning

Build AI Agents with DSPy

In this tutorial, we will walk you through how to build an AI agents with DSPy. AI agents refer to the system that can autonomously perceive its environment, make decisions, and take actions to achieve specific goals. Unlike a single model prompt, an agent typically follows a loop of reasoning, planning, and acting, often integrating tools like search engines, APIs, or memory to complete complex tasks.

This tutorial focuses on a popular architecture of AI agents called **ReAct**, standing for **Reasoning** and **Acting**, which provides a task description along with a list of tools to LM, then lets LM decide whether to call tools for more observations, or generate the final output.

As the demo, let's build a simple airline customer service agent that can do the following:

- Book new trips on behalf of the user.
- Modify existing trips, including flight change and cancellation.
- On tasks it cannot handle, raise a customer support ticket.

We will build it from `dspy.ReAct` module.

Install Dependencies

Before starting, let's install the required packages:

```
!pip install -qu dspy pydantic
```

Define Tools

We need to prepare a list of tools so that the agent can behave like a human airline service agent:

- `fetch_flight_info`: get flight information for certain dates.
- `pick_flight`: pick the best flight based on some criteria.
- `book_flight`: book a flight on behalf of the user.
- `fetch_itinerary`: get the information of a booked itinerary.
- `cancel_itinerary`: cancel a booked itinerary.
- `get_user_info`: get users' information.
- `file_ticket`: file a backlog ticket to have human assist.

Define Data Structure

Before defining the tools, we need to define the data structure. In real production, this will be the database schema. As a demo, we just define the data structure as `pydantic models` for simplicity.

```
In [1]: from pydantic import BaseModel

class Date(BaseModel):
    # Somehow LLM is bad at specifying 'datetime.datetime', so
    # we define a custom class to represent the date.
    year: int
    month: int
    day: int
    hour: int

class UserProfile(BaseModel):
    user_id: str
    name: str
    email: str

class Flight(BaseModel):
    flight_id: str
    date_time: Date
    origin: str
    destination: str
    duration: float
    price: float

class Itinerary(BaseModel):
    confirmation_number: str
    user_profile: UserProfile
    flight: Flight

class Ticket(BaseModel):
    user_request: str
    user_profile: UserProfile
```

Create Dummy Data

Let's also create some dummy data so that the airline agent can do the work. We need to create a few flights and a few users, and initialize empty dictionaries for the itineraries and custom support tickets.

```
In [2]: user_database = {
    "Adam": UserProfile(user_id="1", name="Adam", email="adam@gmail.com"),
    "Bob": UserProfile(user_id="2", name="Bob", email="bob@gmail.com"),
    "Chelsie": UserProfile(user_id="3", name="Chelsie", email="chelsie@gmail.com"),
    "David": UserProfile(user_id="4", name="David", email="david@gmail.com")
}

flight_database = {
    "DA123": Flight(
        flight_id="DA123", # DSPy Airline 123
        origin="SFO",
        destination="JFK",
        date_time=Date(year=2025, month=9, day=1, hour=1),
        duration=3,
        price=200,
    ),
    "DA125": Flight(
        flight_id="DA125",
        origin="SFO",
        destination="JFK",
        date_time=Date(year=2025, month=9, day=1, hour=7),
        duration=9,
        price=500,
    ),
    "DA456": Flight(
        flight_id="DA456",
        origin="SFO",
        destination="SNA",
        date_time=Date(year=2025, month=10, day=1, hour=1),
        duration=2,
        price=100,
    ),
    "DA460": Flight(
        flight_id="DA460",
        origin="SFO",
        destination="SNA",
        date_time=Date(year=2025, month=10, day=1, hour=9),
        duration=2,
        price=120,
    ),
}

itinery_database = {}
ticket_database = {}
```

Define the Tools

Now we can define the tools. In order to have `dspy.ReAct` function properly, every function should:

- Have a docstring which defines what the tool does. If the function name is self-explanable, then you can leave the docstring empty.
- Have type hint for the arguments, which is necessary for LM to generate the arguments in the right format.

```
In [3]: import random
import string

def fetch_flight_info(date: Date, origin: str, destination: str):
    """Fetch flight information from origin to destination on the given date.
    flights = []

    for flight_id, flight in flight_database.items():
        if
            flight.date_time.year == date.year
            and flight.date_time.month == date.month
            and flight.date_time.day == date.day
            and flight.origin == origin
            and flight.destination == destination
        ):
            flights.append(flight)

    if len(flights) == 0:
        raise ValueError("No matching flight found!")
    return flights

def fetch_itinerary(confirmation_number: str):
    """Fetch a booked itinerary information from database"""
    return itinery_database.get(confirmation_number)

def pick_flight(flights: list[Flight]):
    """Pick up the best flight that matches users' request. we pick the
    sorted_flights = sorted(
        flights,
        key=lambda x: (
            x.get("duration") if isinstance(x, dict) else x.duration,
            x.get("price") if isinstance(x, dict) else x.price,
        ),
    )
    return sorted_flights[0]

def _generate_id(length=8):
    chars = string.ascii_lowercase + string.digits
    return "".join(random.choices(chars, k=length))

def book_flight(flight: Flight, user_profile: UserProfile):
    """Book a flight on behalf of the user."""
    confirmation_number = _generate_id()
    while confirmation_number in itinery_database:
        confirmation_number = _generate_id()
    itinery_database[confirmation_number] = Itinerary(
        confirmation_number=confirmation_number,
        user_profile=user_profile,
        flight=flight,
    )
    return confirmation_number, itinery_database[confirmation_number]

def cancel_itinerary(confirmation_number: str, user_profile: UserProfile):
    """Cancel an itinerary on behalf of the user."""
    if confirmation_number in itinery_database:
        del itinery_database[confirmation_number]
    return
    raise ValueError("Cannot find the itinerary, please check your confirmation number")

def get_user_info(name: str):
    """Fetch the user profile from database with given name."""
    return user_database.get(name)

def file_ticket(user_request: str, user_profile: UserProfile):
    """File a customer support ticket if this is something the agent cannot
    ticket_id = _generate_id(length=6)
    ticket_database[ticket_id] = Ticket(
        user_request=user_request,
        user_profile=user_profile,
    )
    return ticket_id
```

Create ReAct Agent

Now we can create the **ReAct** agent via `dspy.ReAct`. We need to provide a signature to `dspy.ReAct` to define task, and the inputs and outputs of the agent, and tell it about the tools it can access.

```
In [4]: import dspy

class DSPyAirlineCustomerService(dspy.Signature):
    """You are an airline customer service agent that helps user book
    flights.

    You are given a list of tools to handle user request, and you should
    fulfill users' request."""
    user_request: str = dspy.InputField()
    process_result: str = dspy.OutputField(
        desc=(
            "Message that summarizes the process result, and the
            confirmation_number if a new flight is booked."
        )
    )

In [5]: agent = dspy.ReAct(
    DSPyAirlineCustomerService,
    tools=[
        fetch_flight_info,
        fetch_itinerary,
        pick_flight,
        book_flight,
        cancel_itinerary,
        get_user_info,
        file_ticket,
    ]
)
```

Use the Agent

To interact with the agent, simply provide the request through `user_request`, and the agent will start doing its job.

Select a language model and set up the API keys. We are using gpt-4o-mini here, but you can change to other models. For how to configure the language model, please refer to this guide.

```
In [6]: import os

os.environ["OPENAI_API_KEY"] = "(your openai key)"
dspy.configure(lm=dspy.LM("openai/gpt-4o-mini"))

In [7]: result = agent(user_request="please help me book a flight from SFO to JFK")
print(result)

{'94ldyhsx': {'confirmation_number': '94ldyhsx', 'user_profile': UserProfile(user_id='1', name='Adam', email='adam@gmail.com'), 'thought_3': "I have gathered Adam's user profile information. Now I can proceed to book the selected flight DA123 for him.", 'tool_name_3': 'book_flight', 'tool_args_3': {'flight': {'flight_id': 'DA123', 'date_time': {'year': 2025, 'month': 9, 'day': 1, 'hour': 1}, 'origin': 'SFO', 'destination': 'JFK', 'duration': 3.0, 'price': 200.0}, 'user_profile': {'user_id': '1', 'name': 'Adam', 'email': 'adam@gmail.com'}}}, 'observation_3': ('94ldyhsx', 'itinerary(confirmation_number='94ldyhsx', user_profile=UserProfile(user_id='1', name='Adam', email='adam@gmail.com'), flight=Flight(flight_id='DA123', date_time=Date(year=2025, month=9, day=1, hour=1), origin='SFO', destination='JFK', duration=3.0, price=200.0))), 'thought_4': "Adam's flight has been successfully booked, and I have received the confirmation number. I will now mark the task as complete.", 'tool_name_4': 'finish', 'tool_args_4': {}, 'observation_4': 'Completed.'}, 'reasoning': "I successfully fetched flight information for the route from SFO to JFK on 09/01/2025. There were two available flights: DA123 and DA125. DA123 is the best option as it is the cheapest and has the shortest duration."}
```

We can see the booked itinerarie in the database.

```
In [8]: print(itinery_database)

{'94ldyhsx': Itinerary(confirmation_number='94ldyhsx', user_profile=UserProfile(user_id='1', name='Adam', email='adam@gmail.com'), flight=Flight(flight_id='DA123', date_time=Date(year=2025, month=9, day=1, hour=1), origin='SFO', destination='JFK', duration=3.0, price=200.0))}
```

Interpret the Result

The result contains the the `process_result` as required by the user, and a `reasoning` field that carries the reasoning behind the answer. In addition, it has a `trajectory` field which contains:

- Reasoning (thought) at each step
- Tools picked by LM at each step
- Arguments for tool calling, determined by LM at each step
- Tool execution results at each step

Behind scene, the `dspy.ReAct` is executing a loop, which accumulates tool call information along with the task description, and send to the LM until hits `max_iters` or the LM decides to wrap up. To better interpret the process, let's use `dspy.LM.inspect_history()` to see what's happening inside each step.

```
In [9]: dspy.LM.inspect_history(n=10)

[2025-05-28T01:06:46.819048]

System message:
Your input fields are:
1. 'user_request' (str)
2. 'trajectory' (str)
Your output fields are:
1. 'next_thought' (str)
2. 'next_tool_name' (Literal['fetch_flight_info', 'fetch_itinerary', 'pick_flight', 'book_flight', 'cancel_itinerary', 'get_user_info', 'file_ticket', 'finish'])
3. 'next_tool_args' (dict[str, Any])
All interactions will be structured in the following way. with the annotation:

We can see that in each LM call, the user message includes the information of previous tool calls, along with the task description.
```

Let's try a different task.

```
In [10]: confirmation_number = "copy the confirmation number here"

result = agent(user_request="I want to take DA125 instead on 09/01, please help me book a flight from SFO to JFK")
print(result)

Prediction(
  trajectory=(
    'thought_0': "I need to fetch the user's current itinerary using the confirmation number provided (7zokt5v5) to understand the details of their existing booking before making any modifications.", 'tool_name_0': 'fetch_itinerary', 'tool_args_0': {'confirmation_number': '7zokt5v5'}, 'observation_0': None, 'thought_1': "I need to fetch the user's current itinerary using the confirmation number provided (7zokt5v5) to understand the details of their existing booking before making any modifications.", 'tool_name_1': 'fetch_itinerary', 'tool_args_1': {'confirmation_number': '7zokt5v5'}, 'observation_1': None, 'thought_2': "I need to file a customer support ticket since I am unable to fetch the user's itinerary with the provided confirmation number.", 'tool_name_2': 'file_ticket', 'tool_args_2': {'user_request': 'I want to take DA125 instead on 09/01, please help me book my itinerary 7zokt5v5', 'user_profile': {'user_id': '1', 'name': 'Adam', 'email': 'adam@gmail.com'}, 'observation_2': 'I need to file a customer support ticket since I have filed a customer support ticket for the user's request. I will wait for a response from the support team.'}
```

Conclusion

Congrats on finishing the tutorial! In this tutorial we have seen how to build a customer service agent with DSPy. The gist are:

- Define the tools as python function, and add **docstring** and **type hints**.
- Provide the tools to `dspy.ReAct` along with a signature to define the task.
- Invoke the `dspy.ReAct` with the inputs field defined in the signature, and it will start the reasoning and acting loop behind the scene.